

COMP3331/9331 Computer Networks and Applications

Assignment for Term 1, 2024

Version 1.0

Due: 11:59am (noon) Thursday, 18 April 2024 (Week 10)

Updates to the assignment, including any corrections and clarifications, will be posted on the course website. Please make sure that you check the course website regularly for updates.

1. Change Log

Version 1.0 released on 7th March 2024.

2. Goal and learning objectives

For this assignment, you are to implement a reliable transport protocol over the UDP protocol. We will refer to the reliable transport protocol that you will be implementing in this assignment as Simple Transport Protocol (STP). STP will include most (but not all) of the features that are described in Sections 3.5.4 - 3.5.6 of the text Computer Networking by Kurose and Ross (7th or 8th ed.) or equivalent parts from the Week 4/5 lecture notes. Examples of these features include timeout, ACK, sequence numbers, sliding window, etc. Note that these features are commonly found in many transport protocols. Therefore, this assignment will give you an opportunity to implement some of these basic features of a transport protocol. In addition, you may have wondered why the designer of the TCP/IP protocol stack includes such a feature-less transport protocol as UDP. You will find in this assignment that you can design your own transport protocol and run it over UDP. This is the case for some multimedia delivery services on the Internet, where they have implemented their own proprietary transport protocol over UDP. QUIC, a newly proposed transport protocol also runs over UDP and implements additional functionalities such as reliability.

Recall that UDP provides point-to-point, unreliable datagram service between a pair of hosts. In this programming assignment, you will develop a more structured protocol, STP, which ensures reliable, end-to-end delivery of data in the face of packet loss. STP provides a byte-stream abstraction like TCP and sends pipelined data segments using a sliding window. However, STP does not implement congestion control or flow control. Finally, whereas TCP allows fully bidirectional communication, your implementation of STP will be asymmetric. There will be two distinct STP endpoints, "sender" and "receiver", respectively. Data packets will only flow in the "forward" direction from the sender to the receiver, while acknowledgments will only flow in the "reverse" direction from the receiver back to the sender. To support reliability in a protocol like STP, state must be maintained at both endpoints. Thus, as in TCP, connection set-up and connection teardown phases will be an integral part of the protocol. STP should implement a sliding window protocol wherein multiple segments can be sent by the sender in a pipelined manner. Like TCP, STP will include some elements of both Go-Back-N (GBN) and Selective Repeat (SR). You will use your STP protocol to transfer a text file (examples provided on the assignment webpage) from the sender to the receiver.

The sender program must also emulate the behaviour of an unreliable communication channel between the sender and receiver. Even though UDP segments can get lost, the likelihood of such losses is virtually zero in our test environment, where the sender and receiver will be executed on the same machine. Further, to properly test the implementation of your sender program, we would like to control the unreliable behaviour of the underlying channel. The sender program should

emulate loss of STP segments in both directions – (i) DATA, SYN, and FIN segments in the forward direction and (ii) ACK segments in the reverse direction. You may assume that the underlying channel will never reorder or corrupt STP segments (in both directions).

Note that it is mandatory that you implement STP over UDP. Do not use TCP sockets. You will not receive any marks for this assignment if you use TCP sockets.

2.1 Learning Objectives

On completing this assignment, you will gain sufficient expertise in the following skills:

1. Detailed understanding of how reliable transport protocols such as TCP function.
2. Socket programming for UDP transport protocol.
3. Protocol and message design.

Non-CSE Student Version: The rationale for this option is that students enrolled in a program that does not include a computer science component have had very limited exposure to programming and in particular working on complex programming assignments. A Non-CSE student is a student who is not enrolled in a CSE program (single or double degree). Examples would include students enrolled exclusively in a **single degree program** such as Mechatronics or Aerospace or Actuarial Studies or Law. **Students enrolled in dual degree programs that include a CSE program as one of the degrees do not qualify.** Any student who meets this criterion and wishes to avail of this option **MUST** email cs3331@cse.unsw.edu.au to seek approval before **5pm, 29th March (Friday, Week 7)**. If approved, we will send you the specification for the non-CSE version of the assignment. We will assume by default that all students are attempting the CSE version of the assignment unless they have sought explicit permission. No exceptions.

3. Assignment Specification

STP should be implemented as two separate programs: sender and receiver. You should implement **unidirectional** transfer of data from the sender to the receiver. As illustrated in Figure 1, data segments will flow from sender to receiver while ACK segments will flow from receiver to sender. The sender and receiver programs will be run from different terminals on the same machine, so you can use localhost, i.e., 127.0.0.1 as the IP address for the sender and receiver in your program. Let us reiterate this, **STP must be implemented on top of UDP**. Do not use TCP sockets. If you use TCP, you will not receive any marks for your assignment.

You will find it useful to review Sections 3.5.4 - 3.5.6 of the text (or the relevant parts from the Week 5 lecture notes). It may also be useful to review the basic concepts of reliable data transfer from Section 3.4 (or relevant parts from the Week 4 lecture notes). Section 3.5 of the textbook which covers the bulk of the discussion on TCP is available to download on the assignment page.




Figure 1: This depicts the assignment setup. A file is to be transferred from the Sender to the Receiver, both running on the same machine. Data segments will flow from the sender to receiver, while ACK segments will flow from the receiver to sender.

3.1 File Names

The main code for the sender should be contained in the following files: `sender.c`, or `Sender.java`, or `sender.py`. You may create additional files such as header files or other class files and name them as you wish.

The sender should accept the following seven arguments:

1. `sender_port`: the UDP port number to be used by the sender to send STP segments to the receiver. The sender will receive ACK segments from the receiver through this same port. We recommend using a random port number between 49152 to 65535 (dynamic port number range) for the sender and receiver ports.
2. `receiver_port`: the UDP port number on which receiver is expecting to receive STP segments from the sender. The receiver should send ACK segments to the sender through this same port. We recommend using a random port number in the same range noted above.
3. `txt_file_to_send`: the name of the text file that must be transferred from sender to receiver using your reliable transport protocol. You may assume that the file included in the argument will be available in the current working directory of the sender with the “read” access permissions set (execute “`chmod +r txt_file_to_send`” at the terminal in the directory containing the file, where `txt_file_to_send` is set to the actual name of the file).
4. `max_win`: the maximum window size in bytes for the sender window. This should be an unsigned integer. Effectively, this is the maximum number of data bytes that the sender can transmit in a pipelined manner and for which ACKs are outstanding. `max_win` must be greater than or equal to 1000 bytes (MSS) and does not include STP headers. When `max_win` is set to 1000 bytes, STP will effectively behave as a stop-and-wait protocol, wherein the sender transmits one data segment at any given time and waits for the corresponding ACK segment. While testing, we will ensure that `max_win` is a multiple of 1000 bytes (e.g., 5000 bytes).
5. `rto`: the value of the retransmission timer in milliseconds. This should be an unsigned integer.
6. `flp`: forward loss probability, which is the probability that any segment in the forward direction (DATA, FIN, SYN) is lost. This should be a float value between 0 and 1 (inclusive). If `flp` is 0.1, then the sender will drop approximately 10% of the segments that it intends to send to the receiver.
7. `rlp`: reverse loss probability, which is the probability of a segment in the reverse direction (i.e., ACKs) being lost. This should be a float value between 0 and 1 (inclusive). If `rlp` is 0.05, then the sender will drop approximately 5% of the ACK segments received.

The sender should be initiated as follows:

If you use Java:

```
java Sender sender_port receiver_port txt_file_to_send max_win rto flp rlp
```

If you use C:

```
./sender sender_port receiver_port txt_file_to_send max_win rto flp rlp
```

If you use Python:

```
python3 sender.py sender_port receiver_port txt_file_to_send max_win rto flp rlp
```

During testing, we will ensure that the 7 arguments provided are in the correct format. We will not test for erroneous arguments, missing arguments, etc. That said, it is good programming practice to check for such input errors.

The main code for the receiver should be contained in the following files: `receiver.c`, or `Receiver.java`, or `receiver.py`. You may create additional files such as header files or other class files and name them as you wish.

The receiver should accept the following four arguments:

1. `receiver_port`: the UDP port number to be used by the receiver to receive STP segments from the sender. This argument should match the second argument for the sender.
2. `sender_port`: the UDP port number to be used by the sender to send STP segments to the receiver. This argument should match the first argument for the sender.
3. `txt_file_received`: the name of the text file into which the data sent by the sender should be stored (this is the file that is being transferred from sender to receiver). You may assume that the receiver program will have permission to create files in its working directory (execute “`chmod +w .`” at the terminal to allow the creation of files in the working directory), and that a file with this name does not exist in the working directory or may be overwritten.
4. `max_win`: the receive window size. This argument should match the fourth argument for the sender. It is provided to the receiver to ensure that the receiver can initialise a buffer sufficient to hold all the data that the sender can transmit in a pipelined manner.

The receiver should be initiated as follows:

If you use Java:

```
java Receiver receiver_port sender_port txt_file_received max_win
```

If you use C:

```
./receiver receiver_port sender_port txt_file_received max_win
```

If you use Python:

```
python3 receiver.py receiver_port sender_port txt_file_received max_win
```

During testing, we will ensure that the 4 arguments provided are in the correct format. We will not test for erroneous arguments, missing arguments, etc. That said, it is good programming practice to check for such input errors.

The receiver must be initiated before initiating the sender. The two programs will be executed on the same machine. Pay attention to the order of the port numbers to be specified in the arguments for the two programs as they are in reverse order (sender port is first for the sender while receiver port is first for the receiver). If you receive an error that one or both port numbers are in use, then choose different values from the dynamic port number range (49152 to 65535) and try again.

The sender and receiver should exit after the file transfer is complete and the required information, as stated in the subsequent sections of this document, is written to the sender and receiver log files.

3.2 Segment Format

STP segments must have 2 *two*-byte fields: “type” and “seqno” headers. Each of these store unsigned integer values in network byte order, i.e., big-endian.



The "type" field takes on 4 possible values. DATA = 0, ACK = 1, SYN = 2, FIN = 3.

Unlike TCP, in which multiple types can be set simultaneously, STP segments must be of exactly one of the types specified above.

The "seqno" field indicates the sequence number of the segment. This field is used in all segments. For DATA segments, the sequence number increases by the size (in bytes) of each segment (excluding headers). For ACK segments, the sequence number acts as a cumulative acknowledgment, and indicates the number of the next byte expected by the receiver. For SYN segments, the sequence number is the initial sequence number (ISN), which should be a randomly chosen integer from 0 to $2^{16} - 1$, which is the maximum sequence number. The sequence number of the first DATA segment of the connection should thus be ISN+1. For FIN segments, the sequence number is one larger than the sequence number of the last byte of the last data segment of the connection. The Maximum Segment Size (MSS) (excluding headers) for STP segment is 1000 bytes. A DATA segment can thus be up to 1004 bytes long. The last DATA segment for the file being transferred may contain less than 1000 bytes as the file size may not be a multiple of 1000 bytes. All segments excluding DATA segments should only contain the headers and must thus be exactly 4 bytes long.

The logic for determining the sequence number and ack number in STP is like TCP. However, STP does not use a separate ack number header field. Rather, the "seqno" field contains the ack number for the ACK segments.

3.3 State Diagram

The asymmetry between sender and receiver leads to somewhat different state diagrams for the two endpoints. The state diagram for STP is shown below, which depicts the normal behaviour for both end points.

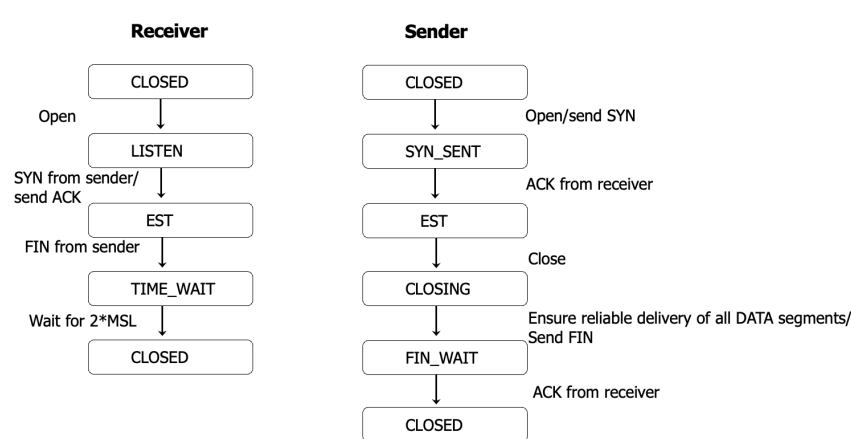


Figure 2: State diagram depicting the normal behaviour of STP.

The receiver can be in four possible states: CLOSED, LISTEN, ESTABLISHED and TIME_WAIT. Initially, it is in the CLOSED state. Upon issuing a passive open, it enters the LISTEN state. Note that the receiver is the passive host in our protocol and is initiated first, while the sender is initiated next and actively opens the connection. While in the LISTEN state, the receiver waits for a SYN segment to arrive on the correct port number. When it does, it responds with an ACK, and moves to the ESTABLISHED state. Should a duplicate SYN arrive while the receiver is in the ESTABLISHED state, it reissues the ACK and remains in the ESTABLISHED state. The ACKs sent by the receiver are cumulative (like TCP). After the sender has reliably transmitted all data

(and received acknowledgments), it will send a FIN segment to the receiver. Upon receipt of the FIN, the receiver moves to the `TIME_WAIT` state. As in TCP, it remains in `TIME_WAIT` for two maximum segment lifetimes (MSLs) before re-entering the `CLOSED` state. This is to ensure that the receiver can respond to potentially retransmitted FIN segments from the sender. You may assume that the MSL is 1 seconds. In other words, the receiver should remain in `TIME_WAIT` for 2 seconds and then transition to `CLOSED`.

The sender can be in five possible states: `CLOSED`, `SYN_SENT`, `ESTABLISHED`, `CLOSING` and `FIN_WAIT`. Like the receiver, the sender starts in the `CLOSED` state. It then issues an active open by sending a SYN segment (to the receiver's port), thus entering the `SYN_SENT` state. This SYN transmission also includes the initial sequence number (ISN) of the conversation. The ISN should be chosen at random from the valid range of possible sequence numbers (0 to $2^{16} - 1$). If a corresponding ACK is not received within `rto` msec, the sender should retransmit the SYN segment. In the common case in which the SYN is acknowledged correctly (the ACK must have the correct sequence number = `ISN + 1`), the sender enters the `ESTABLISHED` state and starts transmitting DATA segments. The sender maintains a single timer (for `rto` msec) for the oldest unacknowledged packet and only retransmits this packet if the timer expires. Once all DATA segments have been transmitted the sender enters the `CLOSING` state. At this point, the sender must still ensure that any buffered data arrives at the receiver reliably. Upon verification of successful transmission, the sender sends a FIN segment with the appropriate sequence number (1 greater than the sequence number of the last data byte) and enters the `FIN_WAIT` state. Once the FIN segment is acknowledged, the sender re-enters the `CLOSED` state. If an ACK is not received before the timer (`rto` msec) expires, the sender should retransmit the FIN segment.

Strictly speaking, you don't have to implement the `CLOSED` state at the start for the sender. Your sender program when executed can immediately send the SYN segment and enter the `SYN_SENT` state. Also, when the sender is in the `FIN_WAIT` state and receives the ACK for the FIN segment, the program can simply exit. This is because the sender only transmits a single file in one execution and quits following the reliable file transfer.

Unlike TCP which follows a three-way handshake (SYN, SYN/ACK, ACK) for connection setup and independent connection closures (FIN, ACK) in each direction, STP follows a two-way connection setup (SYN, ACK) and one directional connection closure (FIN, ACK) process. The setup and closure are always initiated by the sender.

If one end point detects behaviour that is unexpected, it should simply print a message to the terminal indicating that the connection is being reset, and then terminate. For example, if the receiver receives a DATA segment while it is in the `LISTEN` state (where it is expecting a SYN segment). The state transition diagram on the previous page does not capture such erroneous scenarios. You are free to specify the format of the message that is printed. **Note that, we will NOT be rigorously testing your code for such unexpected behaviour.**

3.4 List of features to be implemented by the sender.

You are required to implement the following features in the sender (and equivalent functionality in the receiver).

1. The sender should first open a UDP socket on `sender_port` and initiate a two-way handshake (SYN, ACK) for the connection establishment. The sender sends a SYN segment, and the receiver responds with an ACK. This is different to the three-way handshake implemented by TCP. If the ACK is not received before a timeout (`rto` msec), the sender should retransmit the SYN.
2. The sender must choose a random initial sequence number (ISN) between 0 and $2^{16} - 1$.

Remember to perform sequence number arithmetic modulo 2^{16} . The sequence numbers should cycle back to zero after reaching $2^{16} - 1$.

3. A one-directional (forward) connection termination (FIN, ACK). The sender will initiate the connection close once the entire file has been reliably transmitted by sending the FIN segment and the receiver will respond with an ACK. This is different to the bi-directional close implemented by TCP. If the ACK is not received before a timeout (`rt0` msec), the sender should retransmit the FIN. The sender should terminate after connection closure.

4. STP implements a sliding window protocol like TCP, whereby multiple segments can be transmitted by the sender in a pipelined manner. The sender should maintain a buffer to store all unacknowledged segments. The total amount of data that the sender can transmit in a pipelined manner and for which acknowledgments are pending is limited by `max_win`. Similar to TCP, as the sender receives ACK segments, the left edge of the window can slide forward, and the sender can transmit the next DATA segments (if there is pending data to be sent).

5. Each STP segment transmitted by the sender (including DATA, SYN, FIN) must be encapsulated in a UDP segment and transmitted through the UDP socket.

6. The sender must maintain a single timer for retransmission of data segments (Section 3.5.4 of the text). The value of the timeout will be supplied as a command-line argument to the sender program (`rt0` msec). This timer is for the oldest unacknowledged data segment. In the event of a timeout, **only** the oldest unacknowledged data segment should be retransmitted (like TCP). The sender should not retransmit all unacknowledged segments. Remember that you are NOT implementing Go-Back-N.

7. The sender should implement all the features mentioned in Section 3.5.4 of the text, except for doubling the timeout interval. You are expected to implement the functionality of the simplified TCP sender (Figure 3.33 of the text) and fast retransmit (i.e., the sender should retransmit the oldest unacknowledged data segment on three duplicate ACKs) (pages 247-248).

8. The use of the “seqno” field was outlined in Section 3.2. For DATA segments, the sequence number increases by the size (in bytes) of each segment (excluding headers). For ACK segments, the sequence number acts as a cumulative acknowledgment, and indicates the number of the next byte expected by the receiver. The logic is thus like TCP, except that STP does not use a separate ACK header field. The ACK segments use the seqno header field to indicate the ACK numbers.

9. The sender will receive each ACK segment from the receiver through the same socket which the sender uses to transmit data. The ACK segment will be encapsulated in a UDP segment. The sender must first extract the ACK segment from the UDP segment and then process it as per the operation of the STP protocol. ACK segments have the same format as DATA segments but do not contain any data.

10. The sender program should emulate the behaviour of an unreliable communication channel between the sender and receiver. UDP segments can occasionally experience loss in a network, but the likelihood is very low when the sender and receiver are executed on the same machine. Moreover, to properly test the implementation of your sender program, we would like to control the unreliable behaviour of the underlying channel. The sender emulates loss of STP segments in both directions which can be controlled through two command line argument: (i) `flp`: determines the probability of a DATA, SYN, or FIN segment created by the sender being dropped. In other words, each such segment created by the sender will be dropped with a probability `rlp`. If the segment is not dropped, then it will be transmitted through the socket to the receiver. (ii) `rlp`: determines the probability of an ACK segment in the reverse direction from the sender being dropped. In other words, each ACK segment arriving at the sender socket will be dropped with a probability `flp`. If the segment is not dropped, then it will be processed as per the STP protocol.

Note about Random Number Generation

You will need to generate random numbers to establish an initial sequence number (ISN) and implement segment loss. If you have not learnt about the principles behind random number generators, you need to know that random numbers are in fact generated by a deterministic formula by a computer program. Therefore, strictly speaking, random number generators are called pseudo-random number generators because the numbers are not truly random. The deterministic formula for random number generation in Python, Java and C use an input parameter called a *seed*. If a fixed seed is used, then the same sequence of random numbers will be produced, each time the program is executed. This will thus likely generate the same ISN and the same sequence of segment loss in each execution of the sender. While this may be useful for debugging purposes, it is not a realistic representation of an unreliable channel. Thus, you must ensure that you do not use a fixed seed in your submitted program. A simple way to use a different seed for each execution is to base the seed on the system time.

The following code fragments in Python, Java and C generate random numbers between 0 and 1 with a different seed in each execution.

- In Python, you initialise a random number generator by using `random.seed()`. By default, the random number generator uses the current system time. After that you can generate a random floating point number between (0,1) by using `random.random()`.
- In Java, you initialise a random number generator by using `Random random = new Random()`. This constructor sets the seed of the random number generator to a value very likely to be distinct from any other invocation of this constructor. After that, you can generate a random floating point number between (0,1) by using `float x = random.nextFloat()`.
- In C, you initialise a random number generator by using `srand(time(NULL))`. After that, you can generate a random floating point number between (0,1) by using `float x = rand() / ((float)(RAND_MAX)+1)`. Note that, `RAND_MAX` is the maximum value returned by the `rand()` function.

11. The sender should maintain a log file named `sender_log.txt` where it records the information about each segment that it sends and receives. Information about dropped segments should also be included. You may assume that the sender program will have permission to create files in its current working directory. If the file already exists, then it should be overwritten. Start each entry on a new line. The format should be as follows:

```
<snd|rcv|drp> <time> <type-of-segment> <seq-number> <number-of-bytes>
```

There must be no whitespace within each field, and all fields must be whitespace separated. You may choose how to separate each field, e.g. with padding or a fixed number of spaces/tabs, but the log file must be easily human readable.

The first field is from the perspective of the sender endpoint, i.e., the sender will `snd` DATA, FIN, and SYN segments, if they are not dropped, and `rcv` ACK segments, if they are not dropped. Any dropped segments will be marked as `drp`. The time should be in milliseconds and relative to when the first SYN segment was sent (even if it was dropped), thus the first entry will be a SYN segment at time 0. The type of segment can be ACK, DATA, SYN or FIN. The number of bytes should be zero for all segments other than DATA segments, where it is the number of bytes of actual payload data (i.e., without headers).

The log must be recorded in chronological order.

For example, the following shows the log file for a sender that transmits 3000 bytes of data and the `rto` is 100 msec. The ISN chosen is 4521 and `max_win` is 3000 bytes. Notice that the third DATA segment is dropped and is hence retransmitted after a timeout interval of 100 msec. The interval is quite precise in this example log, but minor variations due to system latency are expected and acceptable.

snd	0.00	SYN	4521	0
rcv	10.34	ACK	4522	0
snd	10.45	DATA	4522	1000
snd	10.55	DATA	5522	1000
drp	10.67	DATA	6522	1000
rcv	36.76	ACK	5522	0
rcv	37.87	ACK	6522	0
snd	137.87	DATA	6522	1000
rcv	167.23	ACK	7522	0
snd	168.11	FIN	7522	0
rcv	203.34	ACK	7523	0

Once the entire file has been reliably transmitted and the connection is closed, the sender should also print the following statistics at the end of the log file (i.e., `sender_log.txt`). Please note that these statistics are specific to the sender while it is in an ESTABLISHED or CLOSING state, i.e., transmitting DATA segments, and receiving the corresponding (and potentially cumulative) ACK segments. We are not including summary statistics related to SYN or FIN segments and their acknowledgement. Also note, the amount of data sent and acknowledged is strictly payload data.

1. Amount of (original) Data Sent (in bytes) – regardless of whether the data is dropped but does not include retransmitted data.
2. Amount of (original) Data Acknowledged (in bytes) – based only on acknowledgements that are not dropped.
3. Number of (original) Data Segments Sent – regardless of whether the segment is dropped but does not include retransmitted segments.
4. Number of Retransmitted Data Segments.
5. Number of Duplicate Acknowledgements Received.
6. Number of Data Segments Dropped.
7. Number of Acknowledgements Dropped.

For the example scenario, the sender log would be appended with:

Original data sent:	3000
Original data acked:	3000
Original segments sent:	3
Retransmitted segments:	1
Dup acks received:	0
Data segments dropped:	1
Ack segments dropped:	0

You must follow the above format, but you may choose how to separate each statistical description with its value, e.g. with padding or a fixed number of spaces/tabs, provided the summary statistics are easily human readable.

NOTE: Generation of this log file is very important. It will help your tutors in understanding the

flow of your implementation and marking. So, if your code does not generate any log files, you will only be graded out of 25% of the marks.

The sender should terminate after the file transfer is complete and the log file is finalised.

The sender must only print a message to the terminal if it is terminating prematurely, due to unexpected behaviour by the receiver. No other output should be displayed. If you are printing output to the terminal for debugging purposes, make sure you disable it prior to submission.

3.5 Specific details about the receiver

1. The receiver should first open a UDP socket on `receiver_port` and then wait for segments to arrive from the sender. The first segment to be sent by the sender is a SYN segment and the receiver will reply with an ACK segment.
2. The receiver should next create a new text file named as per the `txt_file_received` argument. You may assume that the receiver program will have permission to create files in its current working directory (execute “`chmod +w .`” at the terminal to allow the creation of files in the working directory). The received data will be written to this file in the correct order.
3. The receiver should initialise a receive window of `max_win` to hold all data that the sender can send in a pipelined manner.
4. The receiver should generate an ACK immediately after receiving any segment from the sender. The receiver should not follow Table 3.2 of the textbook and does not implement delayed ACKs. The format of the ACK segment is exactly like the STP DATA segment. It should however not contain any data. The ack number should be included in the “seqno” field of the STP segment. There is no explicit ACK field in the STP header.
5. The receiver should buffer all out-of-order data in the receive buffer. This is because STP implements reliable in-order delivery of data.
6. The receiver should write data (in correct order) from the receive buffer to the `txt_file_received` file. The file should be written as correctly ordered data is available to ensure the receiver buffer does not overflow. At the end of the transfer, the receiver should have a duplicate of the text file sent by the sender. You can verify this by using the `diff` command on a Linux machine (e.g., “`diff txt_file_to_send txt_file_received`”). When testing your program, if you have the sender and receiver executing in the same working directory then make sure that the file name provided as the argument to the receiver is different from the file name used by the sender.
7. Once the file transfer is complete, the receiver should follow the state transition process as outlined in Section 3.3 while implementing connection closure. Pay particular attention to the transition from the `TIME_WAIT` state to the `CLOSED` state.
8. The receiver should also maintain a log file named `receiver_log.txt` where it records the information about each segment that it sends and receives, in chronological order. If the file already exists, then it should be overwritten. The format should be exactly similar to the sender log file, as outlined in the sender specification (Section 3.4). The first field is from the perspective of the receiver endpoint, i.e., the receiver will `rcv` DATA, FIN, and SYN segments, and `snd` ACK segments. Time should be in milliseconds and relative to when the first SYN segment is received, thus the first entry will be a SYN segment at time 0. One obvious difference from the sender log is that the receiver log will not record any segments as `drp`.

For example, the following shows the log file for a receiver that matches the scenario outlined in the sender specification (Section 3.4). To recall, the sender transmits 3000 bytes of data and the `rto` is 100 msec. The ISN is 4521, the `max_win` is 3000 bytes, and the third DATA segment is

dropped, triggering a retransmission after the timeout interval.

rcv	0.00	SYN	4521	0
snd	0.34	ACK	4522	0
rcv	10.65	DATA	4522	1000
snd	10.75	ACK	5522	0
rcv	10.95	DATA	5522	1000
snd	11.03	ACK	6522	0
rcv	142.41	DATA	6522	1000
snd	144.58	ACK	7522	0
rcv	170.45	FIN	7522	0
snd	171.59	ACK	7523	0

Once the entire file has been reliably transmitted and the connection is closed (remember to follow the state machine for the receiver), the receiver will also print the following statistics at the end of the log file (i.e., `receiver_log.txt`). Like the sender log, these statistics are specific to the receiver while it is in an ESTABLISHED state, i.e., receiving DATA segments, and sending the corresponding (and potentially cumulative) ACK segments. We are not including summary statistics related to SYN or FIN segments and their acknowledgement. Also note, the amount of data received and acknowledged is strictly payload data.

1. Amount of (original) Data Received (in bytes) – does not include retransmitted data.
2. Number of (original) Data Segments Received – does not include retransmitted data.
3. Number of Duplicate Data Segments Received.
4. Number of Duplicate Acknowledgements Sent.

For the example scenario, the receiver log would be appended with:

Original data received:	3000
Original segments received:	3
Dup data segments received:	0
Dup ack segments sent:	0

You must follow the above format, but you may choose how to separate each statistical description with its value, e.g. with padding or a fixed number of spaces/tabs, provided the summary statistics are easily human readable.

REMINDER: Generation of this log file is very important. It will help your tutors in understanding the flow of your implementation and marking. So, if your code does not generate any log files, you will only be graded out of 25% of the marks.

The receiver should terminate after the file transfer is complete and the log file is finalised.

The receiver must only print a message to the terminal if it is terminating prematurely, due to unexpected behaviour by the sender. No other output should be displayed. If you are printing output to the terminal for debugging purposes, make sure you disable it prior to submission.

3.6 Features Excluded

There are several transport layer features adopted by TCP that are excluded from this assignment:

1. You do not need to implement timeout estimation. The timer value is provided as a command line argument (`rto msec`).
2. You do not need to double the timeout interval.

3. You do not need to implement flow control and congestion control.
4. STP does not have to deal with corrupted or reordered segments. Segments will rarely be corrupted and reordered when the sender and receiver are executing on the same machine. In short, it is safe for you to assume that packets are only lost. Note however, that segments can be dropped by the unreliable channel implemented in the sender program and thus segments may arrive out of order at the receiver.

3.7 Implementation Details

The figure below provides a high-level and simplified view of the assignment. The STP protocol logic implements the state maintained at the sender and receiver, which includes all the state variables and buffers. Note that each STP segment (in each direction) must be encapsulated in a UDP segment and transmitted through the UDP socket at each end point.

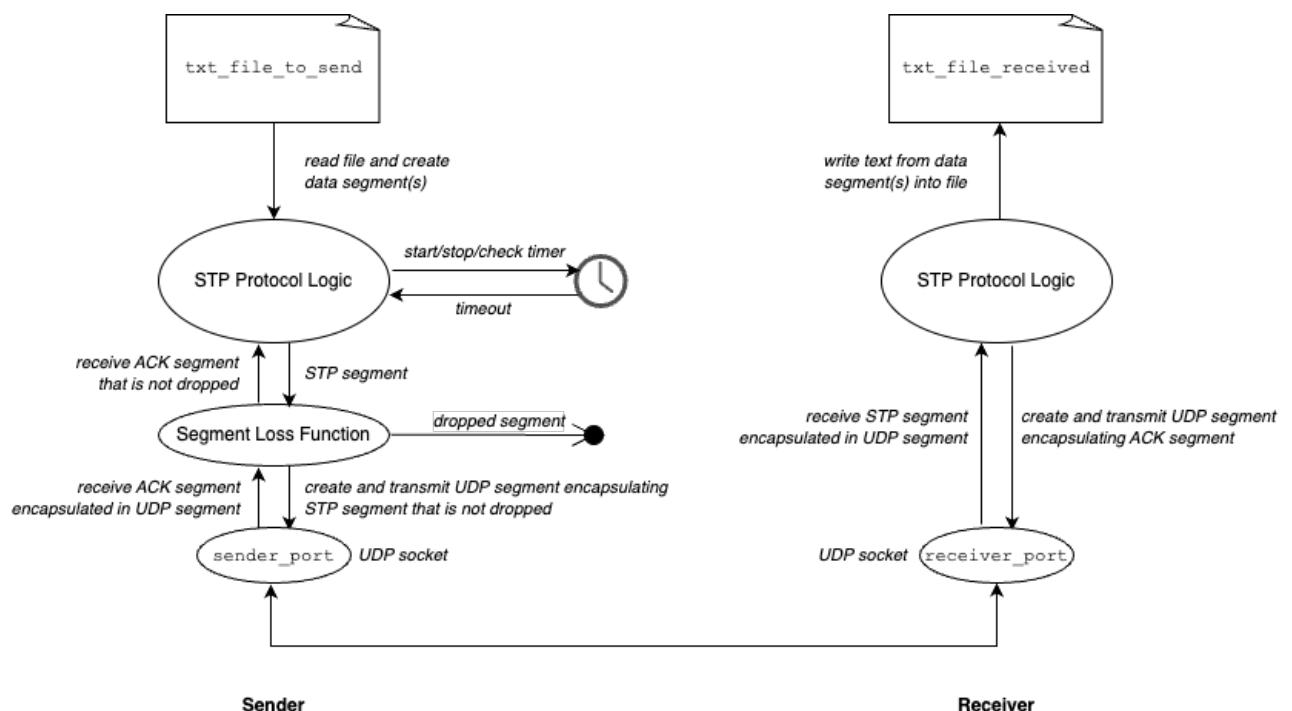


Figure 3: The overall structure of your assignment.

Sender Design

The sender must first execute connection setup, followed by data transmission, and finally connection teardown, as outlined in the state diagram description (Section 3.3). During connection setup, the sender transmits a SYN segment, starts a timer (`rto`) and waits for an ACK. During data transmission, the sender may transmit several STP data segments (determined by `max_win`), all of which need to be buffered (in case of retransmissions) and wait for the corresponding ACKs. A timer (`rto`) should be started for the oldest unacknowledged segment. During connection teardown, the sender must transmit a FIN segment, start a timer (`rto`) and wait for an ACK. Each STP segment (of any kind) that is to be transmitted must be encapsulated in a UDP segment and sent through the sender socket to the receiver. The sender should also process incoming ACK segments from the receiver. Upon receipt of three duplicate ACKs, the sender should retransmit the oldest unacknowledged DATA segment. In the case of a timeout, the sender should retransmit the SYN or FIN, if in the connection setup or teardown process, or retransmit the oldest unacknowledged DATA segment, if in the data transmission process. **As the sender needs to deal with multiple events, we recommend using multi-threading.**

To better understand the logic behind fast retransmission and the use of a single timer, a simplified view of the sender is presented below (this does not include the connection setup and teardown and the functionality for implementing the unreliable channel)

event: data read from text file to be sent

- create STP segment with seqno, where seqno is byte-stream number of first data byte in segment
- transmit segment
- start timer if not already running
 - think of timer as for oldest unACKed segment
 - expiration interval: rto msec

event: timeout

- retransmit oldest unACKed segment
- restart timer
- reset dupACKcount

event: ACK received

- if ACK acknowledges previously unACKed segment(s)
 - update what is known to be ACKed
 - if there are currently any unACKed segments
 - restart timer
 - otherwise (there are no unACKed segments)
 - stop timer
 - reset dupACKcount
- otherwise (ACK acknowledges previously ACKed segments)
 - increment dupACKcount
 - if dupACKcount = 3
 - retransmit oldest unACKed segment
 - reset dupACKcount

Note that the use of a single timer does **not** require the sender to maintain any additional state for each segment such as when it was sent.

Recall also that the sender needs to implement the functionality of an unreliable channel that can drop segments in either direction. Before sending a UDP segment through the socket, the sender should call the function that simulates segment loss with a probability flp . If the segment is to be dropped, then it should not be transmitted and nothing else needs to be done (other than updating the log). Only if the segment is not dropped should it be transmitted through the socket.

Upon receipt of a UDP segment through the socket, the sender should call the function that simulates packet loss with a probability rlp . If the segment is dropped, then nothing else needs to be done (other than updating the log). Remember that we are emulating loss of the incoming segment on the channel, so effectively this segment never arrived at the sender and thus no action is necessary. If the segment is not dropped, only then should the sender execute the STP protocol logic (as outlined above and in Sections 3.3, 3.4).

Receiver Design

Upon receipt of a UDP segment through the socket, the receiver should extract the STP segment which is encapsulated within the UDP segment. The receiver should then execute the STP protocol logic (as outlined in Sections 3.3 and 3.5). For a SYN or FIN segment, the corresponding ACK should be sent, and other actions should be taken as per the state diagram shown in Section 3.3. For a DATA segment, the data should be written to the receive buffer. If the data is received in order,

then it can be written to the file and the buffer space made available. If out of order, then it will remain in the buffer until the missing data is received. An appropriate ACK segment should be generated. The ACK should then be encapsulated in a UDP segment and sent through the receiver socket to the sender.

All the above functionalities can be implemented in a single thread. However, the receiver must implement a timer during connection closure for transitioning from the TIME_WAIT state to the CLOSED state. **The suggested way to implement this is using multi-threading where a separate thread is used to manage the timer.** However, it may be possible to implement this using a single thread and using non-blocking or asynchronous I/O by using polling, i.e., select().

Data Structures

To manage your protocol, you will need a collection of state variables to help you keep track of things like state transitions, sliding windows, and buffers. In TCP, this data structure is referred to as a control block: it's probably a good idea to create a control block class of your own and have a member variable of this type in your primary class. While we do not mandate the specifics, it is critical that you invest some time into thinking about the design of your data structures. You should be particularly careful about how multiple threads will interact with the various data structures.

Note about Artificial Delay

You may be tempted to add artificial delay to your sender and/or receiver through the use of system calls such as `sleep()`. For example, to implement your retransmission timer, or to try and resolve synchronisation issues. You are generally cautioned against doing this. At the very least, it will introduce unnecessary inefficiency to your protocol. Of perhaps greater concern, if you are doing it in response to some issue, then most likely the bug will remain. It will just be more transient and difficult to find. There will almost certainly be a more logical solution that will improve the efficiency and robustness of your programs. In such scenarios, you are highly encouraged to take a systematic approach to debugging your code and properly resolve any issues.

4. Test Implementations

Implementations for both endpoints will be made available on the CSE server system. They will be initiated as follows:

```
3331 sender    sender_port receiver_port txt_file_to_send max_win rto flp rlp
3331 receiver receiver_port sender_port txt_file_received max_win
```

The command-line arguments are as described in Section 3.1.

Please note that the behaviour of your applications is not expected to match the test implementations. For example, the test implementations may save files to a temporary directory, output logs to screen, and provide more explicit details when unexpected behaviour is detected. You should always defer to the assignment specification as to what is expected of your implementations.

Recall, the sender and receiver programs will be run from different terminals on the same machine. For the CSE Linux servers, the easiest way to achieve this is to connect to VLAB and open two terminals. If you are instead connecting over ssh, then you must be mindful that both connections are to the same server. Here are two suggested approaches to using ssh:

1. Start a ssh session and determine the hostname from the shell prompt, e.g., `z1234567@vx01` indicates a connection to host `vx01`. From another local terminal initiate a second ssh session. If you don't connect to the same host then in the second session you can simply ssh to the first host, e.g., `ssh vx01`.

2. Start a ssh session and then use an application like `screen` to multiplex the one terminal between both the sender and receiver applications.

If you receive an error when starting the sender or receiver that a port number is in use, then choose a different value from the dynamic port number range (49152 to 65535) and try again.

5. Additional Notes

- This is NOT a group assignment. You are expected to work on this individually.
- **Sample Code:** We will provide sample code in all 3 languages on the assignment page that you may find useful. You are not required to use it, but you are welcome to adapt it as you see fit to help get started. Sample text files are also provided. We will use different files for our tests.
- **Programming Tutorial:** We will run programming tutorials during regular lab times in Week 7, to get students started with programming some of the building blocks of the assignment. A schedule for these sessions will be announced in Week 6.
- **Assignment Help Sessions:** We will run additional consultations in Weeks 7-10, for all 3 programming languages, to assist you with assignment related questions. A schedule will be posted on the assignment page of the course website. Please note, these sessions are not a forum for tutors to debug your code.
- **Tips on getting started:** The best way to tackle a complex implementation task is to do it in stages. A good starting point is to implement the functionality required for a stop-and-wait protocol (version rdt3.0 from the textbook and lectures), which sends one segment at a time. If you set the `max_win` argument to 1000 bytes (equal to the MSS) for the sender, then it will effectively operate as a stop-and-wait receiver as the sender window can only hold 1 data segment. You can first test with the loss probabilities (`flp`, `rlp`) set to zero to simulate a reliable channel. Once you verify that your protocol works correctly for this setting, you can increase the values for the loss probabilities to test that the sender can work as expected over a channel that loses packets (you may do this progressively, i.e., first only allow for packet loss in the forward direction, then only allow for packet loss in the reverse direction, and finally test with packet loss in both directions). Test comprehensively with different loss probabilities to ensure that your sender works correctly.
- You can next progress to implement the full functionality of STP, wherein the sender should be able to transmit multiple packets in a pipelined manner (i.e., sliding window). First consider the case where the underlying channel is reliable (`flp` and `rlp` are set to 0). Set `max_win` to be a small multiple of the MSS (e.g., 4000 bytes). Once you verify that your protocol works correctly for this setting, you can increase the values for the loss probabilities to test that the sender can work as expected over a channel that loses packets (you may do this progressively, i.e., first only allow for packet loss in the forward direction, then only allow for packet loss in the reverse direction, and finally test with packet loss in both directions). Test comprehensively with different loss probabilities and window sizes to ensure that your sender works correctly.
- You can refer to the following resources for multi-threading. Note that you won't need to implement very complex aspects of multi-threading for this assignment.
 - Python: https://www.tutorialspoint.com/python3/python_multithreading.htm
 - Java: <https://www.javatpoint.com/how-to-create-a-thread-in-java>
 - C: <https://www.geeksforgeeks.org/multithreading-in-c/>
- It is imperative that you rigorously test your code to ensure that all possible (and logical) interactions can be correctly executed. **Test, test, and test.**

- **Debugging:** When implementing a complex assignment such as this, there are bound to be errors in your code. We strongly encourage that you follow a systematic approach to debugging. If you are using an IDE for development, then it is bound to have debugging functionalities. Alternatively, you could use a command line debugger such as pdb (Python), jdb (Java) or gdb (C). Use one of these tools to step through your code, create break points, observe the values of relevant variables and messages exchanged, etc. Proceed step by step, check and eliminate the possible causes until you find the underlying issue. Note that, we won't be able to debug your code on the course forum or even in the help sessions.
- **Backup and Versioning:** We strongly recommend you to back-up your programs frequently. CSE backs up all user accounts nightly. If you are developing code on your personal machine, it is strongly recommended that you undertake daily backups. We also recommend using a good versioning system so that you can roll back and recover from any inadvertent changes. There are many services available for this which are easy to use. If you are using an online versioning system, such as GitHub, then you **MUST** ensure that your repository is private. We will **NOT** entertain any requests for special consideration due to issues related to computer failure, lost files, etc.
- **Language and Platform:** You are free to use C, Java, or Python to implement this assignment. Please choose a language that you are comfortable with. The programs will be tested on CSE Linux machines. So please make sure that your entire application runs correctly in VLAB. This is especially important if you plan to develop and test the programs on your personal computers (which may possibly use a different OS or version or IDE). Note that CSE machines support the following: **gcc version 12.2, Java 17, and Python 3.11**. You may only use the basic socket programming APIs provided in your programming language of choice. You may not use any special ready-to-use libraries or APIs that implement certain functions of the spec for you. If you are unsure, it is best you check with the course staff on the forum.
- You are encouraged to use the course discussion forum to ask questions and to discuss different approaches to solve the problem. However, you should **not** post your solution or any code fragments on the forum.

6. Assignment Submission

Please ensure that you use the mandated file names. You may of course have additional header files and/or helper files. If you are using C, then you **MUST** submit a Makefile along with your code. This is because we need to know how to resolve the dependencies among all the files that you have provided. After running your Makefile we should have the following executable files: `sender` and `receiver`.

In addition, you should submit a small report, `report.pdf` (no more than 2 pages). Provide details of which language you have used (e.g., C) and the organisation of your code (Makefile, directories if any, etc.). Your report must contain a brief discussion of how you have implemented the STP protocol. This should include the overall program design, data structure design, and a brief description of the operation of the sender and receiver. Also discuss any design trade-offs considered and made. If your program does not work under any circumstances, report this here. Also indicate any segments of code that you have borrowed from the Web or other books.

You are required to submit your source code and `report.pdf`. You can submit your assignment using the `give` command through VLAB. Make sure you are in the same directory as your code and report, and then do the following:

1. Type `tar -cvf assign.tar filenames`, e.g.:


```
tar -cvf assign.tar *.java report.pdf
```

2. Next, type: `cs3331 assign assign.tar`

You should receive a message stating the result of your submission, **ensure** it says it's accepted. The same command should be used for 3331 and 9331.

Alternately, you can also submit the tar file via the WebCMS3 interface on the assignment page.

Important notes

- The system will only accept a file named `assign.tar`. All other names will be rejected.
- **Ensure that your program/s are tested in the VLAB environment before submission. In the past, there were cases where tutors were unable to compile and run students' programs while marking. To avoid any disruption, please ensure that you test your program in the VLAB environment before submitting the assignment. Note that, we will be unable to award any significant marks if the submitted code does not run during marking.**
- You may submit as many times as you wish before the deadline. A later submission will override the earlier submission, so make sure you submit the correct file. Do not wait until the last moment to submit, as there may be technical, or network errors, and you will not have time to rectify it.

Late Submission Penalty: The UNSW standard late penalty will apply, which is 5% per day of the maximum available mark, for up to 5 days. This assignment is worth 20 marks, therefore the penalty equates to a 1-mark deduction per day late. Submissions after 5 days will not be accepted.

7. Plagiarism

You are to write all the code for this assignment yourself. All source code is subject to strict checks for plagiarism, via highly sophisticated plagiarism detection software. These checks may include comparison with available code from Internet sites and assignments from previous semesters. In addition, each submission will be checked against all other submissions of the current semester. Do not post this assignment on forums where you can pay programmers to write code for you. We will be monitoring such forums. Please note that we take this matter quite seriously. The LiC will decide on appropriate penalty for detected cases of plagiarism. The most likely penalty would be to reduce the assignment mark to ZERO. We are aware that a lot of learning takes place in student conversations, and don't wish to discourage those. However, it is important, for both those helping others and those being helped, not to provide or accept any programming language code in writing, as this is apt to be used exactly as is, and lead to plagiarism penalties for both the supplier and the copier of the codes. Write something on a piece of paper, by all means, but tear it up/take it away when the discussion is over. It is OK to borrow bits and pieces of code from sample socket code out on the Web and in books. You **MUST** however acknowledge the source of any borrowed code. This means providing a reference to a book or a URL when the code appears (as comments). Also indicate in your report the portions of your code that were borrowed. Explain any modifications you have made (if any) to the borrowed code.

Generative AI Tools: It is prohibited to use any software or service to search for or generate information or answers. If its use is detected, it will be regarded as serious academic misconduct and subject to the standard penalties, which may include 00FL, suspension and exclusion.

8. Marking Policy

You should test your programs rigorously before submitting your code. Your code will be manually marked using the following criteria:

Test 1 - Stop and Wait over a Reliable Channel: 2 marks

We will test your STP implementation when executed as a stop and wait protocol (`max_win=1000`) and when the underlying channel is reliable (`flp=rlp=0`).

We show the instantiation of the two programs assuming the implementation is in Python. The arguments will be similar for C and Java.

```
python3 receiver.py 56007 59606 FileToReceive.txt 1000
python3 sender.py 59606 56007 test1.txt 1000 rto 0 0
```

We will test for different values of `rto` and with different text files. We will compare the received file with the sent file, check the sender and receiver logs, and other checks to ensure that the STP protocol is correctly implemented at both end points.

Test 2 - Stop and Wait over an Unreliable Channel: 4 marks

Next, we will test your STP implementation while operating as a stop and wait protocol (`max_win=1000`) but where the underlying channel is unreliable.

In the first instance, we will only induce packet loss in the forward direction (`rlp=0`). The sender will be instantiated as follows (receiver will be instantiated as above):

```
python3 sender.py 59606 56007 test1.txt 1000 rto flp 0
```

We will test for different values of `rto`, `flp` and with different text files. Checks will be undertaken as noted above. (1 mark)

In the second instance, we will only induce packet loss in the reverse direction (`flp=0`). The sender will be instantiated as follows (receiver will be instantiated as above):

```
python3 sender.py 59606 56007 test1.txt 1000 rto 0 rlp
```

We will test for different values of `rto`, `rlp` and with different text files. Checks will be undertaken as noted above. (1 mark)

In the final instance, we will induce packet loss in both directions. The sender will be instantiated as follows (receiver will be instantiated as above):

```
python3 sender.py 59606 56007 test1.txt 1000 rto flp rlp
```

We will test for different values of `rto`, `flp`, `rlp` and with different text files. Checks will be undertaken as noted above. (2 marks)

Test 3 - Sliding Window over a Reliable Channel: 4 marks

We will test your STP implementation when executed as a sliding window protocol and when the underlying channel is reliable (`flp=rlp=0`).

We show the instantiation of the two programs assuming the implementation is in Python. The arguments will be similar for C and Java.

```
python3 receiver.py 56007 59606 FileToReceive.txt max_win
python3 sender.py 59606 56007 test1.txt max_win rto 0 0
```

We will test for different values of `max_win` (always a multiple of 1000, with the same value provided to both programs), `rto` and with different text files. We will compare the received file with the sent file, check the sender and receiver logs, and other checks to ensure that the STP protocol is correctly implemented at both end points.

Test 4 - Sliding Window over an Unreliable Channel: 6 marks

Next, we will test your STP implementation when executed as a sliding window protocol but where the underlying channel is unreliable.

In the first instance, we will only induce packet loss in the forward direction (`rlp=0`). The sender will be instantiated as follows (receiver will be instantiated as above):

```
python3 sender.py 59606 56007 test1.txt max_win rto flp 0
```

We will test for different values of `max_win`, `rto`, `flp` and with different text files. Checks will be undertaken as noted above. **(2 marks)**

In the second instance, we will only induce packet loss in the reverse direction (`flp=0`). The sender will be instantiated as follows (receiver will be instantiated as above):

```
python3 sender.py 59606 56007 test1.txt max_win rto 0 rlp
```

We will test for different values of `max_win`, `rto`, `rlp` and with different text files. Checks will be undertaken as noted above. **(2 mark)**

In the final instance, we will induce packet loss in both directions. The sender will be instantiated as follows (receiver will be instantiated as above):

```
python3 sender.py 59606 56007 test1.txt max_win rto flp rlp
```

We will test for different values of `max_win`, `rto`, `flp`, `rlp` and with different text files. Checks will be undertaken as noted above. **(2 marks)**

Test 5 - Sliding Window over an Unreliable Channel against Test Implementations: 2 marks

Next, we will test your STP implementations against the test implementations when executed as a sliding window protocol and where the underlying channel is unreliable.

In the first instance, we will test your receiver against our sender. The two programs will be instantiated as follows:

```
python3 receiver.py 56007 59606 FileToReceive.txt max_win  
3331 sender 59606 56007 test1.txt max_win rto flp rlp
```

We will test for different values of `max_win`, `rto`, `flp`, `rlp` and with different text files. Checks will be undertaken as noted above. **(1 mark)**

In the second instance, we will test your sender against our receiver. The two programs will be instantiated as follows:

```
3331 receiver 56007 59606 FileToReceive.txt max_win  
python3 sender.py 59606 56007 test1.txt max_win rto flp rlp
```

We will test for different values of `max_win`, `rto`, `flp`, `rlp` and with different text files. Checks will be undertaken as noted above. **(1 mark)**

Test 6 – Report: 1 mark

The report should not be longer than 2 pages. Provide details of which language you have used (e.g., Python) and the organisation of your code (e.g. Makefiles, directories if any, etc.). Your report must contain a brief discussion of how you have implemented the STP protocol. This should include the overall program design, data structure design, and a brief description of the operation of the sender and receiver. Also discuss any design trade-offs considered and made. If your program does not work under any circumstances, report this here. Also indicate any segments of code that you have borrowed from the Web or other books. We will verify that the description in your report confirms with the actual implementations in the programs.

Test 7 – Properly documented and commented code: 1 mark

We recommend following well-known style guides such as:

Java: <https://google.github.io/styleguide/javaguide.html>

Python: <https://peps.python.org/pep-0008/>

IMPORTANT NOTE: If your sender and receiver do not generate log files as indicated in the specification, you will only be graded out of 25% of the total marks (i.e., a 75% penalty will be assessed).